
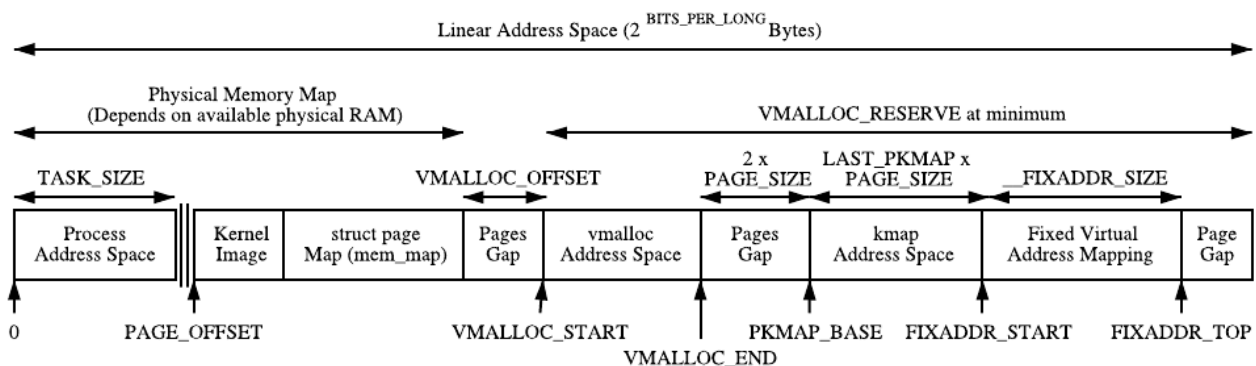



**Figure 1.1 - kernel virtual memory map**

**Note: more during driver architecture**



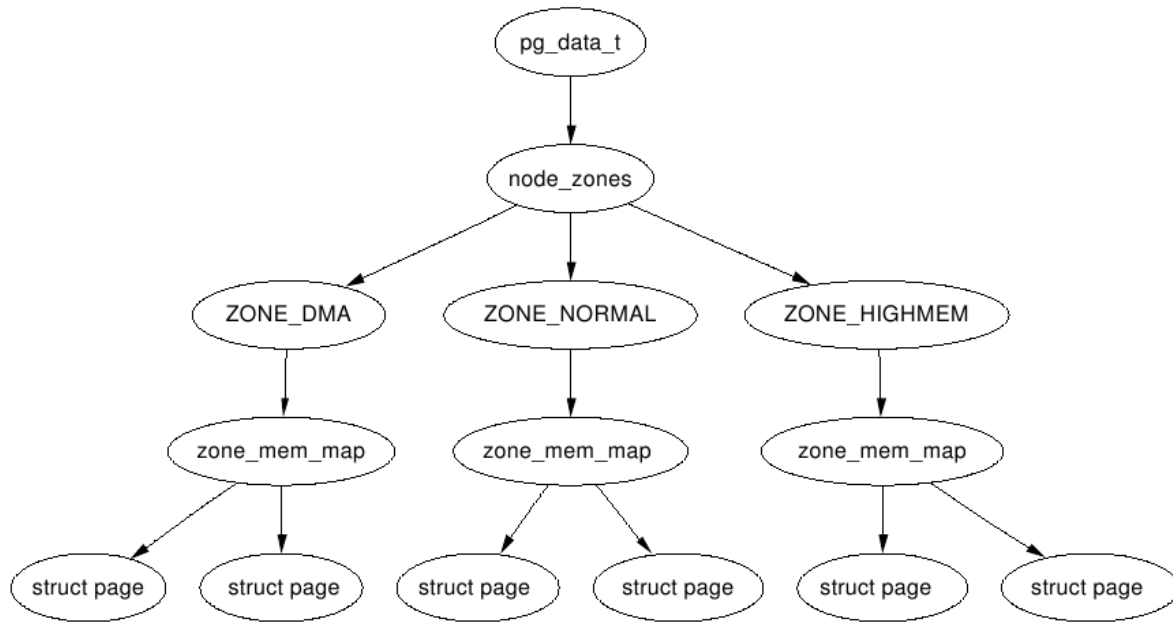
The figure shows the structure of the page table entries used to manage the fourth gigabyte of virtual address space. It indicates the *purpose* of each area of *virtual* address space, and this has nothing to do with the assignment of physical RAM.



**Figure 1.2 - page fault scenarios**

Bit	Set (1)	Not set (0)
0	No page present in RAM	Protection fault (insufficient access permission)
1	Read access	Write access
2	Privileged kernel mode	User mode

**Figure 1.3 - core page-frame management**

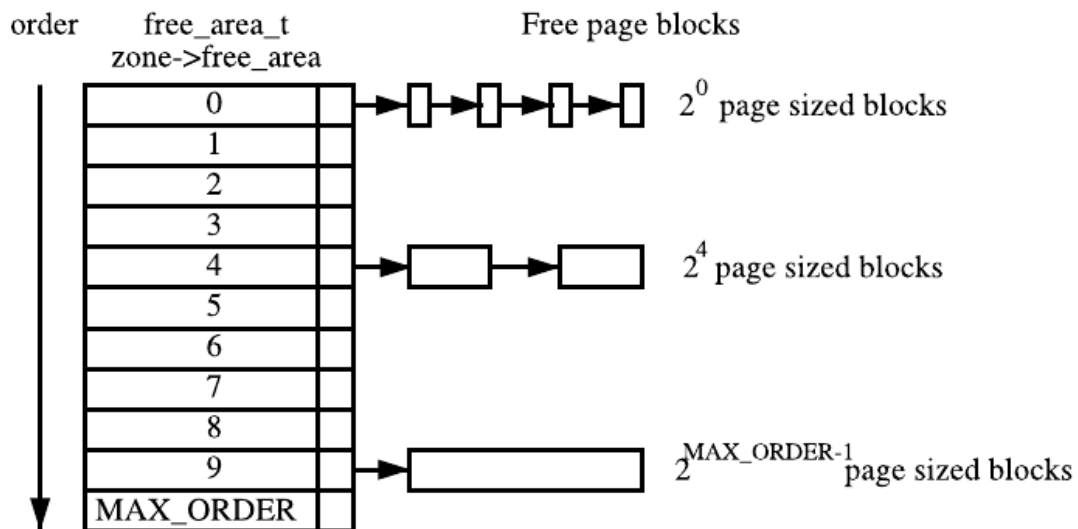


**ZONE\_DMA**            **DMA-able pages**            **< 16MB**

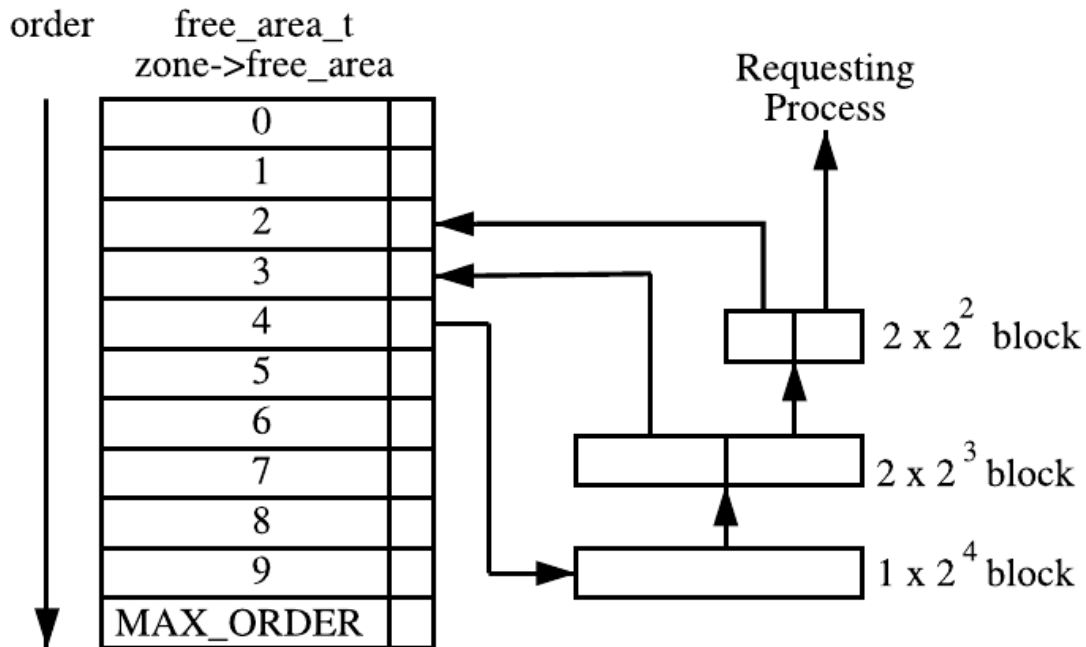
**ZONE\_NORMAL**    **Normally addressable pages** **16–896MB**

**ZONE\_HIGHMEM** **Dynamically mapped pages** **> 896MB**

**Figure 1.4 – buddy allocator per zone**



**Figure 1.5 - buddy allocator in action**



**Figure 1.6 - buddy allocator system APIs**

**struct page \* alloc\_page(unsigned int gfp\_mask)**  
Allocate a single page and return a struct address

**struct page \* alloc\_pages(unsigned int gfp\_mask, unsigned int order)**  
Allocate 2-power-of\_order number of pages and returns a struct page

**unsigned long get\_free\_page(unsigned int gfp\_mask)**  
Allocate a single page, zero it and return a virtual address

**unsigned long \_\_get\_free\_page(unsigned int gfp\_mask)**  
Allocate a single page and return a virtual address

**unsigned long \_\_get\_free\_pages(unsigned int gfp\_mask, unsigned int order)**  
Allocate 2-power-of\_order number of pages and return a virtual address

**void \_\_free\_pages(struct page \*page, unsigned int order)**  
Free an order number of pages from the given page

**void \_\_free\_page(struct page \*page)**  
Free a single page

**void free\_page(void \*addr)**  
Free a page from the given virtual address

**Figure 1.7 - buddy allocator flags**

<b>ZONE_DMA</b>	<b>DMA-able pages</b>	<b>&lt; 16MB</b>
<b>ZONE_NORMAL</b>	<b>Normally addressable pages</b>	<b>16–896MB</b>
<b>ZONE_HIGHMEM</b>	<b>Dynamically mapped pages</b>	<b>&gt; 896MB</b>

**GFP\_KERNEL**            (**\_\_GFP\_WAIT | \_\_GFP\_IO | \_\_GFP\_FS**)

**GFP\_DMA**                **\_\_GFP\_DMA**

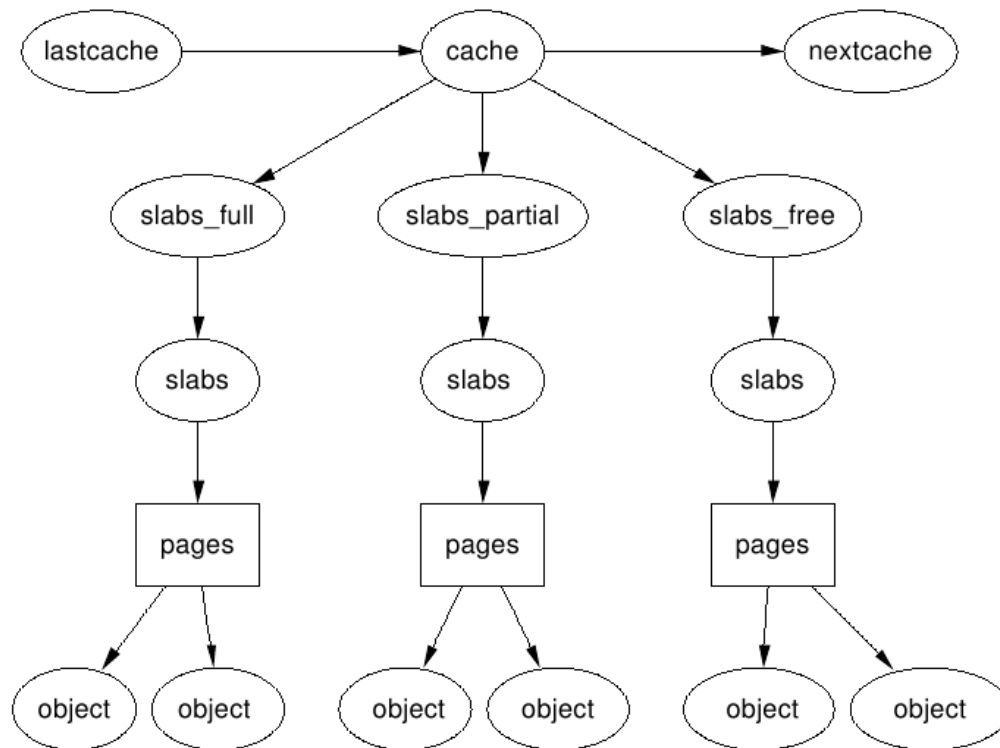
**GFP\_ATOMIC**           **\_\_GFP\_HIGH**

**\_\_GFP\_HIGHMEM**

**Allocates from ZONE\_HIGHMEM or ZONE\_NORMAL**

**Note: more during device drivers**

**Figure 1.8 - slab allocator subsystem architecture**



**Note: we will see more during driver architecture.**